

TP2

April 23, 2019

```
In [215]: # Augustin Commun
          # Chloé Constantineau

import pandas as pd
import xlrd
from IPython.display import display, HTML

df = pd.read_excel('Analyse_du_Credit.xlsx')
display(HTML(df.head().to_html()))
```

<IPython.core.display.HTML object>

```
In [216]: df.describe()
```

```
Out[216]:
```

	credit	loan_duration_months	amount	\
count	1000.000000	1000.000000	1000.000000	
mean	0.300000	20.903000	3271.258000	
std	0.458487	12.058814	2822.736876	
min	0.000000	4.000000	250.000000	
25%	0.000000	12.000000	1365.500000	
50%	0.000000	18.000000	2319.500000	
75%	1.000000	24.000000	3972.250000	
max	1.000000	72.000000	18424.000000	

	payment_percentage_revenu	residence_years	age	bank_credit	\
count	1000.000000	1000.000000	1000.000000	1000.000000	
mean	2.973000	2.845000	35.546000	1.407000	
std	1.118715	1.103718	11.375469	0.577654	
min	1.000000	1.000000	19.000000	1.000000	
25%	2.000000	2.000000	27.000000	1.000000	
50%	3.000000	3.000000	33.000000	1.000000	
75%	4.000000	4.000000	42.000000	2.000000	
max	4.000000	4.000000	75.000000	4.000000	

	dependent
count	1000.000000
mean	1.155000

```

std      0.362086
min      1.000000
25%      1.000000
50%      1.000000
75%      1.000000
max      2.000000

```

```
In [217]: df.groupby('credit').describe()
```

```

Out[217]:      loan_duration_months
           count      mean      std  min  25%  50%  75%
credit
0           700.0  19.207143  11.079564  4.0  12.0  18.0  24.0
1           300.0  24.860000  13.282639  6.0  12.0  24.0  36.0

           amount      ... bank_credit      dependent
           max  count      mean  ...      75%  max      count      mean
credit
0           60.0  700.0  2985.457143  ...      2.0  4.0      700.0  1.155714
1           72.0  300.0  3938.126667  ...      2.0  4.0      300.0  1.153333

           std  min  25%  50%  75%  max
credit
0           0.362844  1.0  1.0  1.0  1.0  2.0
1           0.360911  1.0  1.0  1.0  1.0  2.0

[2 rows x 56 columns]

```

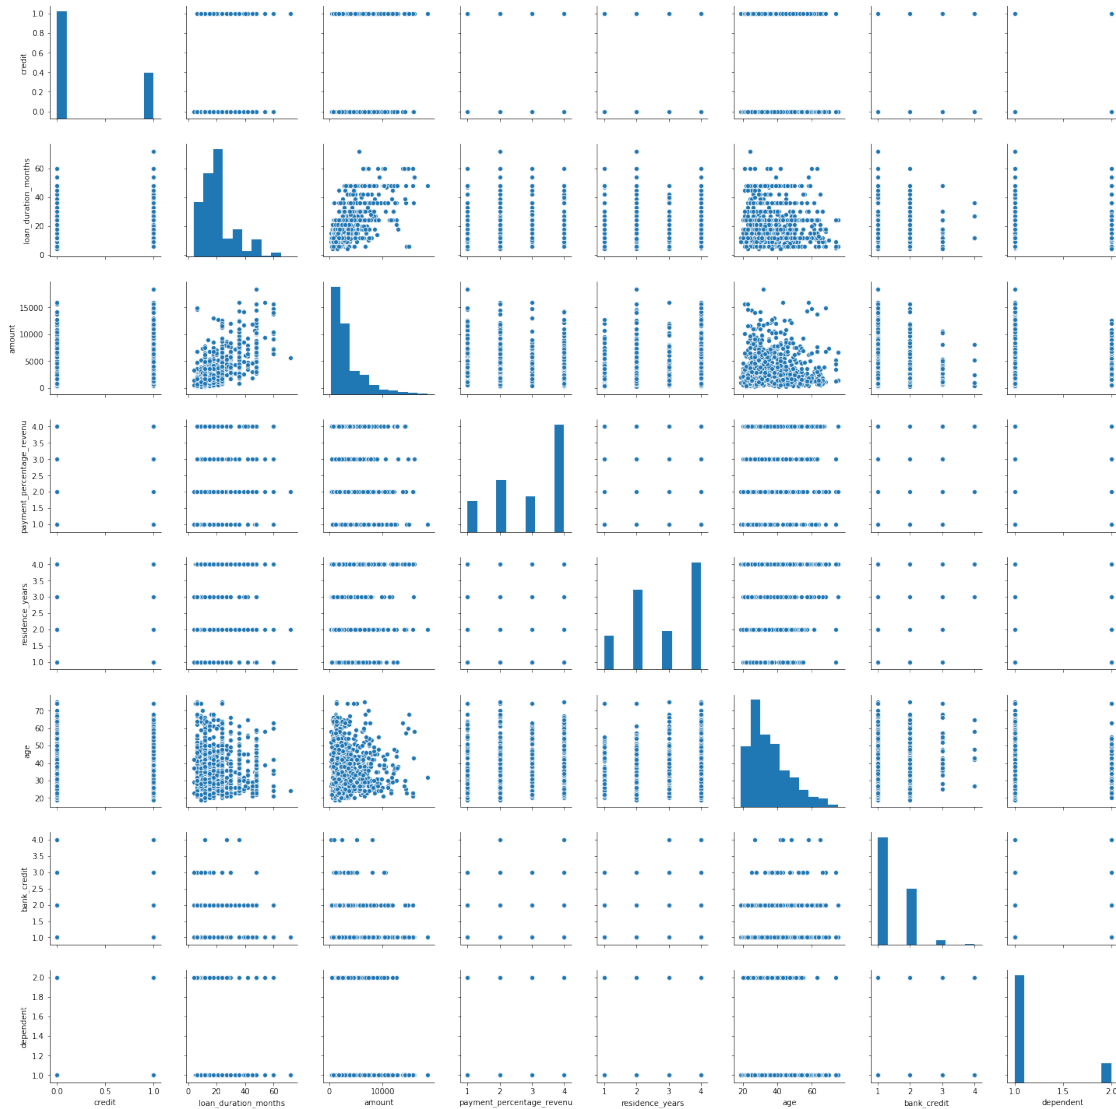
```

In [218]: import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

sns.pairplot(df)


```

```
Out[218]: <seaborn.axisgrid.PairGrid at 0x7fa811c513c8>
```



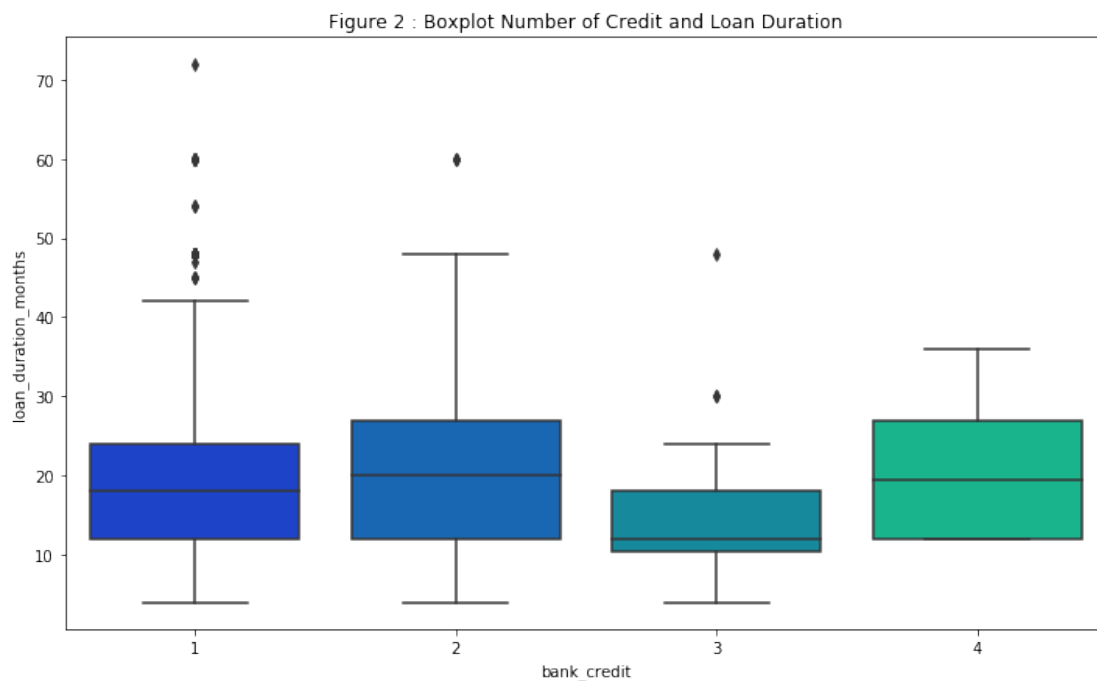
```
In [219]: sns.countplot(x='credit', data=df).set_title("Figure 1 : Credit Count")
```

```
Out[219]: Text(0.5, 1.0, 'Figure 1 : Credit Count')
```



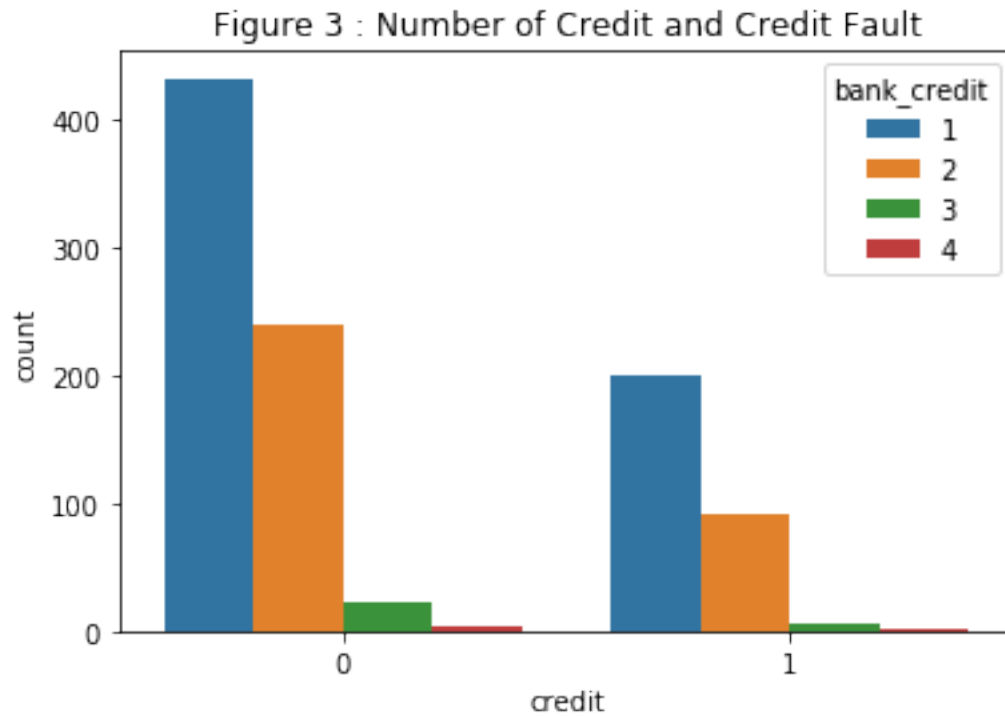
```
In [220]: plt.figure(figsize=(12, 7))
          sns.boxplot(x='bank_credit',y='loan_duration_months',data=df,palette='winter').set_t
```

```
Out[220]: Text(0.5, 1.0, 'Figure 2 : Boxplot Number of Credit and Loan Duration')
```



```
In [221]: sns.countplot(x='credit', hue='bank_credit', data=df).set_title("Figure 3 : Number of Credit and Credit Fault")
```


```
Out[221]: Text(0.5, 1.0, 'Figure 3 : Number of Credit and Credit Fault')
```



```
In [222]: sns.distplot(df['amount']);  
sns.distplot(df['amount'], hist=False, rug=True).set_title("Figure 4 : Dispersion of the Loan Amount")
```

```
Out[222]: Text(0.5, 1.0, 'Figure 4 : Dispersion of the Loan Amount')
```

Figure 4 : Dispersion of the Loan Amount




In [223]: *# Distribution of Credit_Amount for Gender*

```
male_credit = df["amount"].loc[df["sex_marital_status"] != "female : divorced/separated"]
female_credit = df["amount"].loc[df["sex_marital_status"] == "female : divorced/separated"]
total_credit = df['amount'].values
```

```
fig, ax = plt.subplots(1, 3, figsize=(16,4))
```

```
sns.distplot(male_credit, ax=ax[0], color="#FE642E")
ax[0].set_title("Male Credit Amount Distribution", fontsize=16)
sns.distplot(female_credit, ax=ax[1], color="#F781F3")
ax[1].set_title("Female Credit Amount Distribution", fontsize=16)
sns.distplot(total_credit, ax=ax[2], color="#2E64FE")
ax[2].set_title("Total Credit Amount Distribution", fontsize=16)
plt.show()
```



```

In [224]: import plotly.plotly as py
import plotly.figure_factory as ff

import numpy as np
import pandas as pd

bad_credit_amount = df["amount"].loc[df['credit'] == 1].values.tolist()
good_credit_amount = df["amount"].loc[df['credit'] == 0].values.tolist()
bad_duration = df['loan_duration_months'].loc[df['credit'] == 1].values.tolist()
good_duration = df['loan_duration_months'].loc[df['credit'] == 0].values.tolist()

bad_loans = go.Scatter(
    x = bad_duration,
    y = bad_credit_amount,
    name = 'Bad Loans',
    mode = 'markers',
    marker = dict(
        size = 10,
        color = 'rgba(152, 0, 0, .8)',
        line = dict(
            width = 2,
            color = 'rgb(0, 0, 0)'
        )
    )
)

good_loans = go.Scatter(
    x = good_duration,
    y = good_credit_amount,
    name = 'Good Loans',
    mode = 'markers',
    marker = dict(
        size = 10,
        color = 'rgba(34, 139, 34, .9)',
        line = dict(
            width = 2,
        )
    )
)

data = [bad_loans, good_loans]

layout = dict(title = 'Correlation of Risk with <br> Credit Amount Borrowed',
              yaxis = dict(zeroline = False),
              xaxis = dict(zeroline = False))

```

```

    )

    fig = dict(data=data, layout=layout)
    iplot(fig, filename='styled-scatter')

In [99]: import plotly.plotly as py
import plotly.figure_factory as ff

corr = df.corr()

arr_corr = corr.values
arr_corr = np.around(arr_corr, decimals=2)

columns = corr.columns.values.tolist()

fig = ff.create_annotated_heatmap(arr_corr, x=columns, y=columns, colorscale='Reds')

fig.layout.title = 'Correlation'
iplot(fig, filename='annotated_heatmap')

In [225]: import numpy as np
import pandas as pd
from IPython.display import display, HTML

df = pd.read_excel('Analyse_du_Credit.xlsx')

# -- Status_Savings_Account --
condition = df['status_savings_account'] == "< 0 $"
df['status_savings_account_zero'] = np.where(condition, 1, 0)
condition = df['status_savings_account'] == "0 <= ... < 2000 $"
df['status_savings_account_more_zero'] = np.where(condition, 1, 0)

# -- Historical Credit Status --
condition = df['historical_credit_status'] == "critical account/ other credits exist"
df['critical_account'] = np.where(condition, 1, 0)

condition = df['historical_credit_status'] == "existing credits paid back duly till now"
df['existing_credit_paid_back_duly_till_now'] = np.where(condition, 1, 0)

condition = df['historical_credit_status'] == "delay in paying off in the past"
df['delay_in_paying_off'] = np.where(condition, 1, 0)

condition = df['historical_credit_status'] == "all credits at this bank paid back duly till now"
df['all_credit_at_this_bank_paid_back_duly_till_now'] = np.where(condition, 1, 0)

# -- Reason --
condition = df['reason'] == ("furniture/equipment" or "domestic appliances")

```



```

df['reason_domestic'] = np.where(condition, 1, 0)

condition = df['reason'] == "Radio/television"
df['reason_electronics'] = np.where(condition, 1, 0)

condition = df['reason'] == "car(new)"
df['reason_new_car'] = np.where(condition, 1, 0)

condition = df['reason'] == ("car(used)" or "repairs")
df['reason_old_car'] = np.where(condition, 1, 0)

# -- Savings --
condition = df['savings'] == "... < 1000 $"
df['savings_less_1k'] = np.where(condition, 1, 0)

condition = df['savings'] == "1000 <= ... < 5000 $"
df['savings_between_1k_5k'] = np.where(condition, 1, 0)

condition = df['savings'] == "5000 <= ... < 10000 $"
df['savings_between_5k_10k'] = np.where(condition, 1, 0)

condition = df['savings'] == ".. >= 10000 $"
df['savings_more_10k'] = np.where(condition, 1, 0)

# -- Employment Time --
condition = df['employment_time'] == "unemployed"
df['unemployed'] = np.where(condition, 1, 0)

condition = df['employment_time'] == "... < 1 year "
df['employment_less_1_year'] = np.where(condition, 1, 0)

condition = df['employment_time'] == "1 <= ... < 4 years"
df['employment_between_1_4_years'] = np.where(condition, 1, 0)

# -- Sex & Marital Status --
condition = df['sex_marital_status'] == "male : single"
df['male_single'] = np.where(condition, 1, 0)

condition = df['sex_marital_status'] == "male : married/widowed"
df['male_married_widowed'] = np.where(condition, 1, 0)

condition = df['sex_marital_status'] == "male : divorced/separated"
df['male_divorced_seperated'] = np.where(condition, 1, 0)

# -- Other Debtor --
debtor = pd.get_dummies(df['other_debtor'], drop_first=True)
debtor.columns = ['debtor_garantor', 'debtor_none']

```

```

# -- Property type --
df['property_type'] = np.where(df['property_type'].str.contains("real estate"), 1, 0)
property_type = pd.get_dummies(df['property_type'], 'real estate')
property_type.columns = ['property_type_other', 'property_type_real_estate']

# -- other_payment_plan_per_echelon --
other_payment_plan_per_echelon = pd.get_dummies(df['other_payment_plan_per_echelon'])
other_payment_plan_per_echelon.columns = ['other_payment_none', 'other_payment_store']

# -- Dwelling --
dwelling = pd.get_dummies(df['dwelling'], drop_first=True)
dwelling.columns = ['dwelling_owned', 'dwelling_rented']

# -- Job --
condition = df['job'] == "skilled employee / official"
df['job_skilled_employee'] = np.where(condition, 1, 0)

condition = df['job'] == "management/ self-employed/ highly qualified employee/ official"
df['job_management_officer_etc'] = np.where(condition, 1, 0)

# -- Foreign Workers --
foreign_worker = pd.get_dummies(df['foreign_worker'], drop_first=True)
foreign_worker.columns = ['is_foreign_worker']

# Drop useless columns
df.drop(['status_savings_account', 'historical_credit_status', 'reason', 'savings',

# Concat dummy variables
df = pd.concat([df, debtor, property_type, dwelling, foreign_worker],axis=1)

display(HTML(df.head().to_html()))

```

<IPython.core.display.HTML object>

In [226]: df.describe()

```

Out[226]:
          credit  loan_duration_months  amount \
count  1000.000000          1000.000000  1000.000000
mean      0.300000           20.903000   3271.258000
std       0.458487           12.058814   2822.736876
min       0.000000            4.000000    250.000000
25%       0.000000           12.000000   1365.500000
50%       0.000000           18.000000   2319.500000
75%       1.000000           24.000000   3972.250000
max       1.000000           72.000000  18424.000000

          payment_percentage_revenu  residence_years  property_type  age \

```

count	1000.000000	1000.000000	1000.000000	1000.000000
mean	2.973000	2.845000	0.282000	35.546000
std	1.118715	1.103718	0.450198	11.375469
min	1.000000	1.000000	0.000000	19.000000
25%	2.000000	2.000000	0.000000	27.000000
50%	3.000000	3.000000	0.000000	33.000000
75%	4.000000	4.000000	1.000000	42.000000
max	4.000000	4.000000	1.000000	75.000000

	bank_credit	dependent	status_savings_account_zero	...	\
count	1000.000000	1000.000000	1000.000000	...	
mean	1.407000	1.155000	0.274000	...	
std	0.577654	0.362086	0.446232	...	
min	1.000000	1.000000	0.000000	...	
25%	1.000000	1.000000	0.000000	...	
50%	1.000000	1.000000	0.000000	...	
75%	2.000000	1.000000	1.000000	...	
max	4.000000	2.000000	1.000000	...	

	male_divorced_seperated	job_skilled_employee	\
count	1000.000000	1000.000000	
mean	0.050000	0.630000	
std	0.218054	0.483046	
min	0.000000	0.000000	
25%	0.000000	0.000000	
50%	0.000000	1.000000	
75%	0.000000	1.000000	
max	1.000000	1.000000	

	job_management_officer_etc	debtor_garantor	debtor_none	\
count	1000.000000	1000.000000	1000.000000	
mean	0.148000	0.052000	0.907000	
std	0.355278	0.222138	0.290578	
min	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	1.000000	
50%	0.000000	0.000000	1.000000	
75%	0.000000	0.000000	1.000000	
max	1.000000	1.000000	1.000000	

	property_type_other	property_type_real_estate	dwelling_owned	\
count	1000.000000	1000.000000	1000.000000	
mean	0.718000	0.282000	0.713000	
std	0.450198	0.450198	0.452588	
min	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	
50%	1.000000	0.000000	1.000000	
75%	1.000000	1.000000	1.000000	
max	1.000000	1.000000	1.000000	

	dwelling_rented	is_foreign_worker
count	1000.000000	1000.000000
mean	0.179000	0.963000
std	0.383544	0.188856
min	0.000000	0.000000
25%	0.000000	1.000000
50%	0.000000	1.000000
75%	0.000000	1.000000
max	1.000000	1.000000

[8 rows x 38 columns]

In [227]: `df.groupby('credit').describe()`

```
Out[227]:
```

	loan_duration_months							
	count	mean	std	min	25%	50%	75%	
credit								
0	700.0	19.207143	11.079564	4.0	12.0	18.0	24.0	
1	300.0	24.860000	13.282639	6.0	12.0	24.0	36.0	

	amount		... dwelling_rented		is_foreign_worker		
	max	count	mean	...	75%	max	count
credit				...			
0	60.0	700.0	2985.457143	...	0.0	1.0	700.0
1	72.0	300.0	3938.126667	...	0.0	1.0	300.0

	mean	std	min	25%	50%	75%	max
credit							
0	0.952857	0.212096	0.0	1.0	1.0	1.0	1.0
1	0.986667	0.114889	0.0	1.0	1.0	1.0	1.0

[2 rows x 296 columns]

In [228]: *# First try with all variables*

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score,
import math

X_train, X_test, y_train, y_test = train_test_split(df.drop('credit',axis=1),
                                                    df['credit'], test_size=0.30,
                                                    random_state=42)

logit_model = LogisticRegression()
logit_model.fit(X_train,y_train)
```

```

predictions = logit_model.predict(X_test)
print(classification_report(y_test,predictions))

# Calcul du AIC avec toutes les variables
resid = y_test - predictions
sse = sum(resid**2)
k = len(df.columns) - 1
AIC= 2*k - 2*math.log(sse)
print("AIC")
print(AIC)

# Score
score = logit_model.score(X_test, y_test)

# Confusion Matrix
confusion_matrix = confusion_matrix(y_test, predictions)
plt.figure(figsize=(9,9))
sns.heatmap(confusion_matrix, annot=True, fmt=".3f", linewidths=.5, square = True, c
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
all_sample_title = 'Confusion Matrix -- Accuracy Score: {0}'.format(score)
plt.title(all_sample_title, size = 15);

# ROC -- TEST data
logit_roc_auc = roc_auc_score(y_test, logit_model.predict(X_test))
fpr, tpr, thresholds = roc_curve(y_test, logit_model.predict_proba(X_test)[: ,1])
plt.figure()
plt.plot(fpr, tpr, label='Logistic Regression -- TEST data (area = %0.2f)' % logit_
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()

# ROC -- TRAIN data
logit_roc_auc = roc_auc_score(y_train, logit_model.predict(X_train))
fpr, tpr, thresholds = roc_curve(y_train, logit_model.predict_proba(X_train)[: ,1])
plt.figure()
plt.plot(fpr, tpr, label='Logistic Regression -- TRAIN data (area = %0.2f)' % logit_
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')

```

```
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()
```


/home/chloe/.local/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning:



Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

	precision	recall	f1-score	support
0	0.77	0.89	0.83	209
1	0.62	0.40	0.48	91
micro avg	0.74	0.74	0.74	300
macro avg	0.70	0.65	0.66	300
weighted avg	0.73	0.74	0.72	300

AIC

65.31238915629262





In [229]: *# SMOTE*

```
from imblearn.over_sampling import SMOTE

data_final = df.copy()

X = data_final.loc[:, data_final.columns != 'credit']
y = data_final.loc[:, data_final.columns == 'credit']

os = SMOTE(random_state=0)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)

columns = X_train.columns
os_data_X, os_data_y = os.fit_sample(X_train, y_train)
os_data_X = pd.DataFrame(data=os_data_X, columns=columns)
os_data_y = pd.DataFrame(data=os_data_y, columns=['credit'])

# we can check the numbers of our data
print("length of oversampled data is ", len(os_data_X))
print("Number of no subscription in oversampled data", len(os_data_y[os_data_y['credit'] == 0]))
print("Number of subscription", len(os_data_y[os_data_y['credit'] == 1]))
print("Proportion of no subscription data in oversampled data is ", len(os_data_y[os_data_y['credit'] == 0]) / len(os_data_X))
print("Proportion of subscription data in oversampled data is ", len(os_data_y[os_data_y['credit'] == 1]) / len(os_data_X))
```

length of oversampled data is 972

Number of no subscription in oversampled data 486

Number of subscription 486

Proportion of no subscription data in oversampled data is 0.5

Proportion of subscription data in oversampled data is 0.5

/home/chloe/.local/lib/python3.6/site-packages/sklearn/utils/validation.py:761: DataConversionWarning:

A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n,) or (1, n).

```
In [230]: from sklearn.feature_selection import RFE
          from sklearn.linear_model import LogisticRegression
```

```
data_final_vars = data_final.columns.values.tolist()
y = ['credit']
X = [i for i in data_final_vars if i not in y]

logreg = LogisticRegression()
rfe = RFE(logreg, 20)
rfe = rfe.fit(os_data_X, os_data_y.values.ravel())
```

```

column_ranking = np.asarray(rfe.ranking_)
get_important_indexes = lambda x, xs: [i for (y, i) in zip(xs, range(len(xs))) if x =
number_ones = get_indexes(1,myarray)
important_column_names = []

print("The most important variables are : ")

for i in number_ones:
    important_column_names.append(df.columns[i])

print(important_column_names)

```

The most important variables are :

```
['residence_years', 'dependent', 'status_savings_account_zero', 'status_savings_account_more_z
```

```
/home/chloe/.local/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarn
```

Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

```
/home/chloe/.local/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarn
```

Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

```
/home/chloe/.local/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarn
```

Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

```
/home/chloe/.local/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarn
```

Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

```
/home/chloe/.local/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarn
```

Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

```
/home/chloe/.local/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarn
```

Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

```
/home/chloe/.local/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarn
```

Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

```
/home/chloe/.local/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarn
```

Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

Deuxième essai avec les variables retenues par l'élimination récursive

```
df_copy = df[['residence_years', 'dependent', 'status_savings_account_zero',
               'status_savings_account_more_zero', 'critical_account',
               'existing_credit_paid_back_duly_till_now', 'delay_in_paying_off',
               'savings_more_10k', 'employment_between_1_4_years', 'male_single',
               'male_married_widowed', 'male_divorced_seperated', 'job_skilled_employ',
               'job_management_officer_etc', 'debtor_garantor', 'debtor_none',
               'property_type_other', 'property_type_real_estate', 'dwelling_owned',
               'dwelling_rented', 'credit']].copy();
```

```
X_train, X_test, y_train, y_test = train_test_split(df_copy.drop('credit',axis=1),
                                                    df_copy['credit'], test_size=0.3,
                                                    random_state=42)
```

```
logit_model = LogisticRegression()
logit_model.fit(X_train,y_train)
predictions = logit_model.predict(X_test)
print(classification_report(y_test,predictions))
```

Calcul du AIC avec toutes les variables

```
resid = y_test - predictions
sse = sum(resid**2)
k = len(df.columns) - 1
AIC= 2*k - 2*math.log(sse)
print("AIC")
print(AIC)
```

Score

```
score = logit_model.score(X_test, y_test)
```

Confusion Matrix

```
confusion_matrix = confusion_matrix(y_test, predictions)
plt.figure(figsize=(9,9))
sns.heatmap(confusion_matrix, annot=True, fmt=".3f", linewidths=.5, square = True, c
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
all_sample_title = 'Confusion Matrix -- Accuracy Score: {0}'.format(score)
plt.title(all_sample_title, size = 15);
```

ROC -- TEST data

```
logit_roc_auc = roc_auc_score(y_test, logit_model.predict(X_test))
fpr, tpr, thresholds = roc_curve(y_test, logit_model.predict_proba(X_test)[:,:1])
plt.figure()
plt.plot(fpr, tpr, label='Logistic Regression -- TEST data (area = %0.2f)' % logit_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
```

```

plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()

# ROC -- TRAIN data
logit_roc_auc = roc_auc_score(y_train, logit_model.predict(X_train))
fpr, tpr, thresholds = roc_curve(y_train, logit_model.predict_proba(X_train)[: ,1])
plt.figure()
plt.plot(fpr, tpr, label='Logistic Regression -- TRAIN data (area = %0.2f)' % logit_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
plt.show()

```


/home/chloe/.local/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning:



Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

	precision	recall	f1-score	support
0	0.75	0.93	0.83	209
1	0.63	0.29	0.39	91
micro avg	0.73	0.73	0.73	300
macro avg	0.69	0.61	0.61	300
weighted avg	0.71	0.73	0.70	300

AIC

65.23594673065224





1 Explication des meilleures variables

L'algorithme d'élimination récursive semble avoir trouvé des variables pertinentes pour les modèles.

En observant le AIC des deux modèles (toutes les variables vs. variables sélectionnées par l'algorithme d'élimination récursive) on remarque que celui avec les variables sélectionnées a le plus petit AIC. Ainsi, il semble être le modèle à privilégier. Par contre, en comparant le score (accuracy) des deux modèles, on remarque que celui avec toutes les variables possède une meilleure performance avec une marge de 1%.

Malgré que la marge de différence est très faible, les résultats sont un peu contradictoire par rapport à la norme qui est que le modèle avec le plus faible AIC est à privilégier. Cependant, nous allons continuer à explorer ces deux modèles avec d'autres fonctions de classification pour voir ce qu'il en est et comment les deux modèles évoluent.

```
In [232]: # Decision Tree - toutes les variables
```

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
from sklearn import tree
import math

X_train, X_test, y_train, y_test = train_test_split(df.drop('credit',axis=1),
                                                    df['credit'], test_size=0.30,
                                                    random_state=42)

tree_gini = DecisionTreeClassifier(criterion = "gini", random_state = 42,
                                   max_depth=10, min_samples_leaf=15)
tree_gini.fit(X_train, y_train)

predictions = tree_gini.predict(X_test)

print(classification_report(y_test,predictions))
print(predictions)

# Score
score = tree_gini.score(X_test, y_test)

# Confusion Matrix
confusion_matrix = confusion_matrix(y_test, predictions)
plt.figure(figsize=(9,9))
sns.heatmap(confusion_matrix, annot=True, fmt=".3f", linewidths=.5, square = True, c
plt.ylabel('Actual label');
```



```

plt.xlabel('Predicted label');
all_sample_title = 'Confusion Matrix -- Accuracy Score: {0}'.format(score)
plt.title(all_sample_title, size = 15);

# ROC -- TEST data
tree_gini_roc_auc = roc_auc_score(y_test, tree_gini.predict(X_test))
fpr, tpr, thresholds = roc_curve(y_test, tree_gini.predict_proba(X_test)[:,-1])
plt.figure()
plt.plot(fpr, tpr, label='Decision Tree Gini - TEST DATA (area = %0.2f)' % tree_gini_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Tree_ROC')
plt.show()

# ROC -- TRAIN data
tree_gini_roc_auc = roc_auc_score(y_train, tree_gini.predict(X_train))
fpr, tpr, thresholds = roc_curve(y_train, tree_gini.predict_proba(X_train)[:,-1])
plt.figure()
plt.plot(fpr, tpr, label='Decision Tree Gini - TRAIN DATA (area = %0.2f)' % tree_gini_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Tree_ROC')
plt.show()

```

	precision	recall	f1-score	support
0	0.77	0.89	0.83	209
1	0.61	0.40	0.48	91
micro avg	0.74	0.74	0.74	300
macro avg	0.69	0.64	0.65	300
weighted avg	0.72	0.74	0.72	300

```


[0 0 0 0 0 0 0 0 0 0 0 0 1 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 1
0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 1 1 0 0 0 1 1 0 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0
0 1 0 0 0 1 0 0 0 0 0 0 1 1 1 0 0 0 0 1 0 0 0 0 0 1 0 0 1 0 0 0 0 0 1 0 0 1
0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 1 1 0 1 0 0 1 0 0 0 0 1 0 1 0 0 0
1 0 0 1 0 0 0 0 0 0 0 0 1 1 0 1 1 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 0 1 0 0 0 0 0



```

```

0 1 0 0 0 0 0 0 0 0 0 1 0 0 0 0 1 0 0 0 0 0 1 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 1 0 0 1 0 0 0 0 0 1 0 1 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 1
0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 1 0 0 0 1 1 0 1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0
1 0 1 0]

```





```

In [233]: # Decision Tree - avec quelques variables clé

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
from sklearn import tree
import math

few_variables = df[['residence_years', 'dependent', 'status_savings_account_zero',
                    'status_savings_account_more_zero', 'critical_account',
                    'existing_credit_paid_back_duly_till_now', 'delay_in_paying_off',
                    'savings_more_10k', 'employment_between_1_4_years', 'male_single',
                    'male_married_widowed', 'male_divorced_seperated', 'job_skilled_employ',
                    'job_management_officer_etc', 'debtor_garantor', 'debtor_none',
                    'property_type_other', 'property_type_real_estate', 'dwelling_owned',
                    'dwelling_rented', 'credit']].copy()

X_train, X_test, y_train, y_test = train_test_split(few_variables.drop('credit',axis=1),
                                                    few_variables['credit'], test_size=0.3,
                                                    random_state=42)

tree_gini = DecisionTreeClassifier(criterion = "gini", random_state = 42,
                                   max_depth=10, min_samples_leaf=15)
tree_gini.fit(X_train, y_train)

predictions = tree_gini.predict(X_test)

print(classification_report(y_test,predictions))
print(predictions)

# Score
score = tree_gini.score(X_test, y_test)

# Confusion Matrix
confusion_matrix = confusion_matrix(y_test, predictions)
plt.figure(figsize=(9,9))
sns.heatmap(confusion_matrix, annot=True, fmt=".3f", linewidths=.5, square = True, cbar=True)
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
all_sample_title = 'Confusion Matrix -- Accuracy Score: {0}'.format(score)
plt.title(all_sample_title, size = 15);

```

```

# ROC -- Test
tree_gini_roc_auc = roc_auc_score(y_test, tree_gini.predict(X_test))
fpr, tpr, thresholds = roc_curve(y_test, tree_gini.predict_proba(X_test)[: ,1])
plt.figure()
plt.plot(fpr, tpr, label='Decision Tree Gini - TEST DATA (area = %0.2f)' % tree_gini_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Tree_ROC')
plt.show()

```

```

# ROC -- Train
tree_gini_roc_auc = roc_auc_score(y_train, tree_gini.predict(X_train))
fpr, tpr, thresholds = roc_curve(y_train, tree_gini.predict_proba(X_train)[: ,1])
plt.figure()
plt.plot(fpr, tpr, label='Decision Tree Gini - TRAIN DATA (area = %0.2f)' % tree_gini_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Tree_ROC')
plt.show()

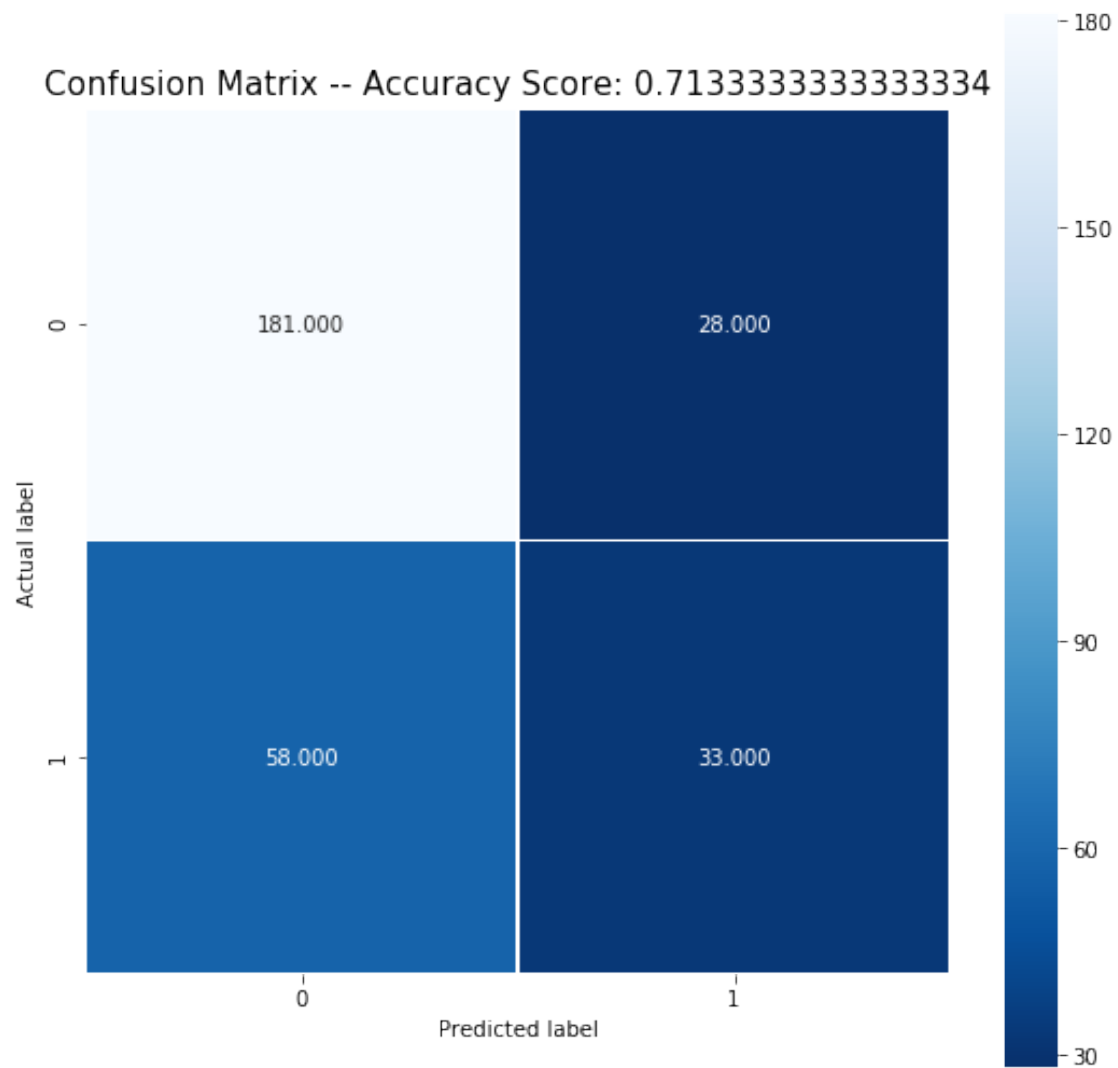
```


	precision	recall	f1-score	support
0	0.76	0.87	0.81	209
1	0.54	0.36	0.43	91
micro avg	0.71	0.71	0.71	300
macro avg	0.65	0.61	0.62	300
weighted avg	0.69	0.71	0.69	300

```

[0 1 1 0 0 0 0 0 0 0 0 1 0 0 1 0 0 0 0 1 0 1 0 0 0 1 1 0 0 0 0 0 0 0 0 1 0
0 0 0 0 0 0 0 1 1 1 0 0 0 0 0 1 0 0 0 0 1 0 0 0 0 1 0 0 0 0 1 0 0 1 0 0 0
0 0 0 0 0 1 0 0 0 0 0 0 0 1 0 0 0 0 1 0 0 0 0 1 1 0 0 1 0 0 0 0 0 1 0 1 0
0 0 0 0 0 0 0 0 0 0 1 0 0 1 0 1 0 0 0 0 0 1 0 0 0 0 1 0 0 0 0 1 0 0 1 0 0
1 1 0 1 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 1 0 1 0 0 0 1 0 0 0 0 0 1 0 0 0 0 1
0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 1 0 0 0 0 0 0 1 1 0 0 0 1 0 0 0 0 0 1 0 0 0
0 0 0 0 0 0 0 1 0 0 0 0 0 1 0 0 1 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 1 0 0 0
0 0 0 0 0 0 0 0 1 1 1 0 0 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 0
0 0 1 0]

```





2 EXPLICATION DES DIVERGENCES

Après avoir essayer plusieurs variations de valeur de paramètre du modèle decision tree, nous avons conclu que d'avoir le `min_samples_leaf` à 15 et le `max_depth` à 10 nous donnait les meilleurs résultats. Nous avons gardé le `random_state` de toutes les variables (incluant le modèle Logit) à 42 par soucis de constance.

2.0.1 Comparaison variables - Logit vs Decision Tree avec Gini

Dans le cas du decision tree, les variables trouvées par l'algorithme d'élimination récursive n'obtiennent toujours pas un meilleur résultat que le modèle avec toutes les variables. Par contre, le grand changement est que son ROC avec les training data est à 70, ce qui est un très bon score. Ainsi, ce modèle a possiblement beaucoup de potentiel par rapport aux autres. Il faudrait tester avec plus de data pour conclure davantage.

2.0.2 Classifier à privilégier

À ce stade ci, le classifier qui a eu le plus de succès est le modèle logit avec toutes les variables. C'est ce modèle qui a obtenu le plus haut score d'accuracy (74.3%) en plus d'avoir le deuxièmes meilleur ROC (0.67 avec les variables de training). Il est toutefois à noter que certains autres essais ont obtenu des résultats similaires (0.61, 0.64, etc.).

2.1 Divergence des deux modèles

Il se peut que le modèle logit fonctionne mieux puisqu'il s'agit d'un problème de séparation linéaire. Ce type de problème/données obtient généralement de bons résultats avec un classifieur Logit.

In [234]: *# Random Forest -- toutes les variables*

```
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
from sklearn import tree
import math

X_train, X_test, y_train, y_test = train_test_split(df.drop('credit',axis=1),
                                                    df['credit'], test_size=0.30,
                                                    random_state=42)

randomForest = RandomForestClassifier(n_estimators=200, max_depth=200,
                                     random_state=42)

randomForest.fit(X_train, y_train)
```



```

predictions = randomForest.predict(X_test)

print(classification_report(y_test,predictions))
print(predictions)

# Score
score = randomForest.score(X_test, y_test)

# Confusion Matrix
confusion_matrix = confusion_matrix(y_test, predictions)
plt.figure(figsize=(9,9))
sns.heatmap(confusion_matrix, annot=True, fmt=".3f", linewidths=.5, square = True, c
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
all_sample_title = 'Confusion Matrix -- Accuracy Score: {0}'.format(score)
plt.title(all_sample_title, size = 15);

# ROC -- Test
random_forest_gini_roc_auc = roc_auc_score(y_test, randomForest.predict(X_test))
fpr, tpr, thresholds = roc_curve(y_test, randomForest.predict_proba(X_test)[: ,1])
plt.figure()
plt.plot(fpr, tpr, label='Random Forest Gini - TEST DATA (area = %0.2f)' % random_fo
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Tree_ROC')
plt.show()

# ROC -- Train
random_forest_gini_roc_auc = roc_auc_score(y_train, randomForest.predict(X_train))
fpr, tpr, thresholds = roc_curve(y_train, randomForest.predict_proba(X_train)[: ,1])
plt.figure()
plt.plot(fpr, tpr, label='Random Forest Gini - TRAIN DATA (area = %0.2f)' % random_f
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Tree_ROC')
plt.show()

importances = randomForest.feature_importances_

```

```

std = np.std([tree.feature_importances_ for tree in randomForest.estimators_],
              axis=0)

indices = np.argsort(importances)[::-1]
x_axis_labels = []
features = list(df.columns.values)
features.pop(0)

# Print the feature ranking
print("Feature ranking:")
for f in range(X_test.shape[1]):
    print("%d. feature %s (%f)" % (f + 1, features[indices[f]], importances[indices[f]]))
    x_axis_labels.append(features[indices[f]])

# Plot the feature importances of the forest
plt.figure(figsize=[15, 5])
plt.title("Feature importances")
plt.bar(range(X_test.shape[1]), importances[indices],
        color="r", yerr=std[indices], align="center")
plt.xticks(range(X_test.shape[1]), x_axis_labels, rotation=90)
plt.xlim([-1, X_test.shape[1]])
plt.show()


```

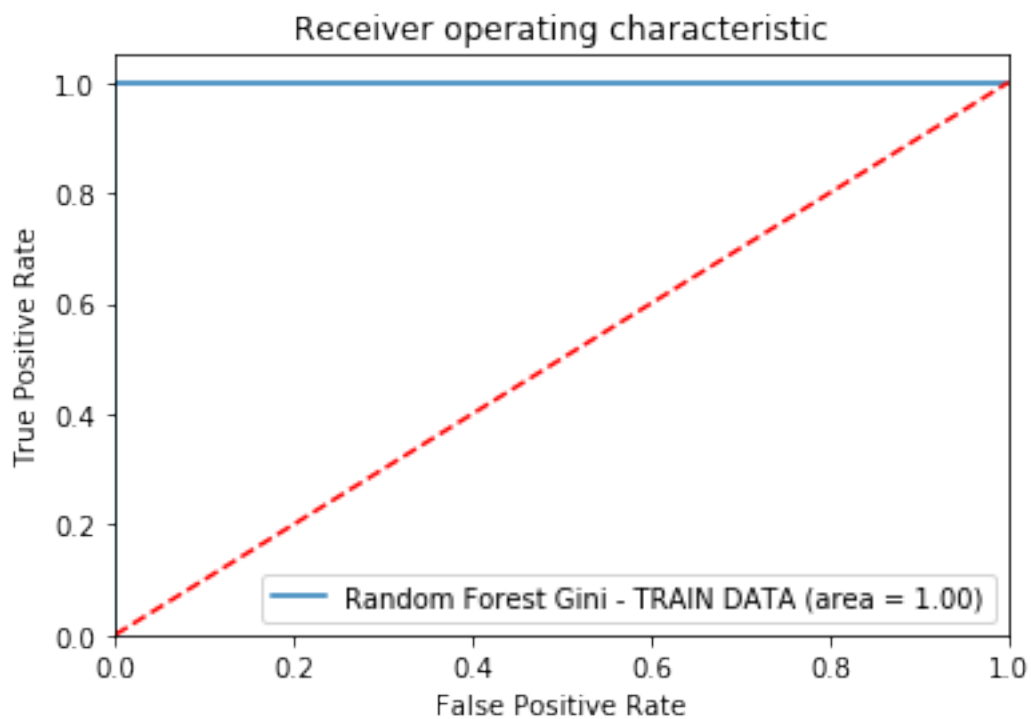

	precision	recall	f1-score	support
0	0.77	0.93	0.84	209
1	0.70	0.36	0.48	91
micro avg	0.76	0.76	0.76	300
macro avg	0.74	0.65	0.66	300
weighted avg	0.75	0.76	0.73	300

```

[0 0 1 0 0 0 0 0 0 0 0 1 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0
 0 0 1 0 0 0 0 0 1 1 0 0 0 1 0 0 0 0 0 0 0 1 1 0 0 1 1 0 0 0 0 1 0 0 0 0 0 0
 0 1 0 0 0 1 0 0 0 0 0 0 1 1 0 0 0 0 1 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 1
 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 1 0 0 1 0 0
 1 1 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 1 0 1 0 1 0 0 0 1 1 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0
 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 1
 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 1 0 1 0 0 0
 0 0 1 0]


```





Feature ranking:

1. feature amount (0.154390)
2. feature age (0.124560)
3. feature loan_duration_months (0.106424)
4. feature status_savings_account_zero (0.059120)
5. feature payment_percentage_revenu (0.051558)
6. feature residence_years (0.049067)
7. feature critical_account (0.033673)
8. feature status_savings_account_more_zero (0.032052)
9. feature savings_less_1k (0.030683)
10. feature bank_credit (0.029361)
11. feature male_single (0.026589)
12. feature job_skilled_employee (0.024970)
13. feature dwelling_owned (0.024950)
14. feature employment_between_1_4_years (0.024644)
15. feature emplyment_less_1_year (0.021567)
16. feature existing_credit_paid_back_duly_till_now (0.020169)
17. feature job_management_officer_etc (0.017198)
18. feature debtor_none (0.016251)
19. feature dependent (0.015689)
20. feature all_credit_at_this_bank_paid_back_duly_till_now (0.015242)
21. feature unemployed (0.014551)
22. feature dwelling_rented (0.014091)
23. feature delay_in_paying_off (0.011951)
24. feature property_type_other (0.011852)
25. feature male_married_widowed (0.011376)
26. feature property_type (0.011308)
27. feature property_type_real_estate (0.011150)
28. feature savings_between_1k_5k (0.010593)
29. feature debtor_garantor (0.008548)
30. feature male_divorced_seperated (0.008526)
31. feature is_foreign_worker (0.004464)
32. feature reason_domestic (0.003432)
33. feature reason_electronics (0.000000)
34. feature reason_new_car (0.000000)
35. feature reason_old_car (0.000000)
36. feature savings_between_5k_10k (0.000000)
37. feature savings_more_10k (0.000000)



In [235]: *# Random Forest -- avec quelques variables clé*

```
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
from sklearn import tree
from sklearn.base import clone
import math
```

```
few_variables = df[['residence_years', 'dependent', 'status_savings_account_zero',
                    'status_savings_account_more_zero', 'critical_account',
                    'existing_credit_paid_back_duly_till_now', 'delay_in_paying_off',
                    'savings_more_10k', 'employment_between_1_4_years', 'male_single',
                    'male_married_widowed', 'male_divorced_seperated', 'job_skilled_employee',
                    'job_management_officer_etc', 'debtor_garantor', 'debtor_none',
                    'property_type_other', 'property_type_real_estate', 'dwelling_owned',
                    'dwelling_rented', 'credit']].copy()
```

```
X_train, X_test, y_train, y_test = train_test_split(few_variables.drop('credit', axis=1),
                                                    few_variables['credit'], test_size=0.2,
                                                    random_state=42)
```

```

randomForest = RandomForestClassifier(n_estimators=200, max_depth=200,
                                     random_state=0)

randomForest.fit(X_train, y_train)

predictions = randomForest.predict(X_test)

print(classification_report(y_test, predictions))
print(predictions)

# Score
score = randomForest.score(X_test, y_test)

# Confusion Matrix
confusion_matrix = confusion_matrix(y_test, predictions)
plt.figure(figsize=(9,9))
sns.heatmap(confusion_matrix, annot=True, fmt=".3f", linewidths=.5, square = True, c
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
all_sample_title = 'Confusion Matrix -- Accuracy Score: {0}'.format(score)
plt.title(all_sample_title, size = 15);

# ROC -- Test
random_forest_gini_roc_auc = roc_auc_score(y_test, randomForest.predict(X_test))
fpr, tpr, thresholds = roc_curve(y_test, randomForest.predict_proba(X_test)[: ,1])
plt.figure()
plt.plot(fpr, tpr, label='Random Forest Gini - TEST DATA (area = %0.2f)' % random_fo
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Tree_ROC')
plt.show()

# ROC -- Train
random_forest_gini_roc_auc = roc_auc_score(y_train, randomForest.predict(X_train))
fpr, tpr, thresholds = roc_curve(y_train, randomForest.predict_proba(X_train)[: ,1])
plt.figure()
plt.plot(fpr, tpr, label='Random Forest Gini - TRAIN DATA (area = %0.2f)' % random_f
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')

```

```

plt.legend(loc="lower right")
plt.savefig('Tree_ROC')
plt.show()

features = ['residence_years', 'dependent', 'status_savings_account_zero',
            'status_savings_account_more_zero', 'critical_account',
            'existing_credit_paid_back_duly_till_now', 'delay_in_paying_off',
            'savings_more_10k', 'employment_between_1_4_years', 'male_single',
            'male_married_widowed', 'male_divorced_seperated', 'job_skilled_employee',
            'job_management_officer_etc', 'debtor_garantor', 'debtor_none',
            'property_type_other', 'property_type_real_estate', 'dwelling_owned',
            'dwelling_rented', 'credit']

importances = randomForest.feature_importances_
std = np.std([tree.feature_importances_ for tree in randomForest.estimators_],
              axis=0)

indices = np.argsort(importances)[::-1]
x_axis_labels = []
# Print the feature ranking
print("Feature ranking:")
for f in range(X_test.shape[1]):
    print("%d. feature %s (%f)" % (f + 1, features[indices[f]], importances[indices[f]]))
    x_axis_labels.append(features[indices[f]])

# Plot the feature importances of the forest
plt.figure(figsize=[19, 5])
plt.title("Feature importances")
plt.bar(range(X_test.shape[1]), importances[indices],
        color="r", yerr=std[indices], align="center")
plt.xticks(range(X_test.shape[1]), x_axis_labels, rotation=90)
plt.xlim([-1, X_test.shape[1]])
plt.show()

```

	precision	recall	f1-score	support
0	0.75	0.80	0.78	209
1	0.46	0.38	0.42	91
micro avg	0.68	0.68	0.68	300
macro avg	0.61	0.59	0.60	300
weighted avg	0.66	0.68	0.67	300

```

[0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 0 0 0 0 1 0 1 0 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 1 1 0 0 0 0 0 0 0 1 0 0 0 0 1 1 0 0 1 1 0 0 1 0 1 1 0 1 0 0 1
0 0 0 0 1 1 0 0 0 1 0 0 1 1 0 0 0 0 1 1 0 0 0 0 0 0 0 0 1 0 0 0 0 1 0 0 1 1
0 0 0 0 0 1 0 0 0 0 1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 1 0 0 1 0 0
1 1 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 1 0 0 0 0 0 1 0 0 1 0 0 1 0 0 0 0 0
0 0 0 0 0 0 1 0 1 0 0 1 1 1 0 1 0 1 0 1 0 0 1 0 0 0 1 0 0 1 0 0 0 1 0 1 0


```





```

0 0 0 0 0 0 0 1 0 0 1 0 0 1 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 1 1 0 0 0
0 0 0 0 0 0 0 0 1 1 1 0 0 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 1 1 1 1 1 0 0 0
0 0 1 1]


```





Feature ranking:

1. feature residence_years (0.170161)
2. feature status_savings_account_zero (0.094704)
3. feature status_savings_account_more_zero (0.075401)
4. feature employment_between_1_4_years (0.073280)
5. feature male_single (0.063604)
6. feature job_skilled_employee (0.061062)
7. feature critical_account (0.060890)
8. feature existing_credit_paid_back_duly_till_now (0.056026)
9. feature dwelling_owned (0.046240)
10. feature job_management_officer_etc (0.040835)
11. feature dependent (0.040316)
12. feature debtor_none (0.038044)
13. feature delay_in_paying_off (0.033574)
14. feature male_married_widowed (0.029370)
15. feature dwelling_rented (0.028339)
16. feature property_type_real_estate (0.025026)
17. feature property_type_other (0.024847)
18. feature male_divorced_seperated (0.022129)
19. feature debtor_garantor (0.016151)
20. feature savings_more_10k (0.000000)



2.2 Explication des meilleures variables

Les résultats indiquent qu'encore une fois, le meilleur modèle est celui avec toutes les variables et non celui avec les variables éliminé récursivement. Cependant, on remarque que sur les 37 variables du modèle, il y en a 5 qui ne sont pas utiles (feature reason_electronics, feature reason_new_car, feature reason_old_car, feature savings_between_5k_10k, feature savings_more_10k). Cette dernière se retrouve aussi avec un score de 0.00 dans notre second modèle à moins de variable.

En regardant ces 5 variables avec aucun impact, on remarque qu'il s'agit de catégorie de 'raison d'un prêt'. Ainsi, la raison d'un prêt ne semble pas avoir d'importance sur les risques de défaut de paiement du crédit. Les autres variables font parti de la catégorie 'Montant d'épargne'. Il s'agit de catégorie à montants substantiels. Ainsi, un individu avec beaucoup d'épargne est moins susceptible d'être en défaut de paiement, ce qui semble bien évident et donc nos résultats semblent cohérents.

On remarque que l'importance des variables entre les 2 modèles sont similaire, mais n'apparaissent pas dans le même ordre. Par exemple, `residence_years` est le facteur numéro 1 dans notre modèle à moins de variable, mais se retrouve en 6e position dans notre modèle avec toutes les variables.

Un fait important à noter avec ce classificateur est qu'il est celui avec un des pires score d'accuracy, mais il s'agit du meilleur classificateur en terme de courbe ROC du training data. Alors, il est possible qu'avec plus de data, le modèle soit en mesure de mieux classer le test data.

3 Sélection de variable pour les prochains algorithmes

Pour les prochaines sections, nous avons choisi de retenir les 20 premières variables du modèle Random Forest testé avec toutes les variables. Il s'agit donc des 20 variables les plus pertinentes sur un total de 37.

1. feature amount (0.151034)
2. feature age (0.120030)
3. feature loan_duration_months (0.108575)
4. feature status_savings_account_zero (0.064480)
5. feature payment_percentage_revenu (0.052127)
6. feature residence_years (0.049769)
7. feature status_savings_account_more_zero (0.034736)
8. feature savings_less_1k (0.032910)
9. feature bank_credit (0.028540)
10. feature critical_account (0.027082)
11. feature job_skilled_employee (0.026977)
12. feature male_single (0.025685)
13. feature employment_between_1_4_years (0.022066)
14. feature dwelling_owned (0.020568)
15. feature employment_less_1_year (0.019408)
16. feature existing_credit_paid_back_duly_till_now (0.019130)
17. feature all_credit_at_this_bank_paid_back_duly_till_now (0.019000)
18. feature job_management_officer_etc (0.018594)
19. feature dwelling_rented (0.016747)
20. feature dependent (0.016100)

In [236]: *## SVM Linéaire avec quelques variables*

```
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
from sklearn.metrics import accuracy_score
```

```

from sklearn.svm import LinearSVC
from sklearn.calibration import CalibratedClassifierCV
import math

few_variables = df[['credit', 'amount', 'age', 'loan_duration_months', 'status_saving',
'payment_percentage_revenu', 'residence_years', 'status_savings_account_more_zero',
'bank_credit', 'critical_account', 'job_skilled_employee', 'male_single',
'employment_between_1_4_years', 'dwelling_owned', 'employment_less_1_year',
'existing_credit_paid_back_duly_till_now', 'all_credit_at_this_bank_paid_back_duly_till_now',
'job_management_officer_etc', 'dwelling_rented', 'dependent']].copy()

X_train, X_test, y_train, y_test = train_test_split(few_variables.drop('credit',axis=1),
                                                    few_variables['credit'], test_size=0.2,
                                                    random_state=42)

clf = LinearSVC(dual=False, max_iter=90000)
clf.fit(X_train, y_train)
predictions = clf.predict(X_test)

print(classification_report(y_test,predictions))
print(predictions)

# Score
score = clf.score(X_test, y_test)

# Confusion Matrix
confusion_matrix = confusion_matrix(y_test, predictions)
plt.figure(figsize=(9,9))
sns.heatmap(confusion_matrix, annot=True, fmt=".3f", linewidths=.5, square = True, cmap='magma')
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
all_sample_title = 'Confusion Matrix -- Accuracy Score: {0}'.format(score)
plt.title(all_sample_title, size = 15);

```

	precision	recall	f1-score	support
0	0.77	0.90	0.83	209
1	0.62	0.36	0.46	91
micro avg	0.74	0.74	0.74	300
macro avg	0.69	0.63	0.64	300
weighted avg	0.72	0.74	0.72	300

```


[0 0 0 0 0 0 0 1 0 0 0 1 0 1 1 0 0 0 1 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 1
0 0 1 0 0 0 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 1 0 0 1 0 0 0 0 0 1 0 0 0 0 0 1
1 0 0 0 0 0 0 0 0 0 0 0 1 1 1 0 0 0 0 1 1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 1
0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 1 0 0 0 1 0 1 0 0 1 0 0
0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0

```

```

0 0 0 0 0 0 1 0 0 0 0 1 1 1 0 1 0 0 0 0 1 0 0 0 0 0 1 1 0 1 0 0 0 0 0 0
0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 1
0 1 0 0 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 1 0 1 0 1 0 0 0
0 0 0 1]

```



In [237]: *## SVM Radial avec quelques variables*

```

from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
from sklearn.metrics import accuracy_score
from sklearn.svm import SVC

```

```

from sklearn.calibration import CalibratedClassifierCV
import math

few_variables = df[['credit', 'amount', 'age', 'loan_duration_months', 'status_saving',
                    'payment_percentage_revenu', 'residence_years', 'status_savings_account_more_zero',
                    'bank_credit', 'critical_account', 'job_skilled_employee', 'male_single',
                    'employment_between_1_4_years', 'dwelling_owned', 'emplyment_less_1_year',
                    'existing_credit_paid_back_duly_till_now', 'all_credit_at_this_bank_paid_back_duly_till_now',
                    'job_management_officer_etc', 'dwelling_rented', 'dependent']].copy()

X_train, X_test, y_train, y_test = train_test_split(few_variables.drop('credit',axis=1),
                                                    few_variables['credit'], test_size=0.2,
                                                    random_state=42)

clf = SVC(C=1.0, kernel='rbf',
          degree=3, gamma='scale',
          coef0=0.0, shrinking=True, probability=False,
          tol=0.001, cache_size=200, class_weight=None,
          verbose=False, max_iter=-1, decision_function_shape='ovr',
          random_state=42)

clf.fit(X_train, y_train)
predictions = clf.predict(X_test)

# Score
score = clf.score(X_test, y_test)


print(classification_report(y_test,predictions))
print(predictions)

# Confusion Matrix
confusion_matrix = confusion_matrix(y_test, predictions)
plt.figure(figsize=(9,9))
sns.heatmap(confusion_matrix, annot=True, fmt=".3f", linewidths=.5, square = True, cmap='Blues')
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
all_sample_title = 'Confusion Matrix -- Accuracy Score: {0}'.format(score)
plt.title(all_sample_title, size = 15);

```

	precision	recall	f1-score	support
0	0.71	1.00	0.83	209
1	1.00	0.05	0.10	91
micro avg	0.71	0.71	0.71	300
macro avg	0.85	0.53	0.47	300
weighted avg	0.80	0.71	0.61	300

```
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 1 0  
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
0 0 0 0]
```



In [238]: *## Réseau de Neurone avec quelques variables*

```
from sklearn.model_selection import train_test_split
```



```

from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
from sklearn.metrics import accuracy_score
from sklearn.neural_network import MLPClassifier
import math

few_variables = df[['credit', 'amount', 'age', 'loan_duration_months', 'status_saving',
    'payment_percentage_revenu', 'residence_years', 'status_savings_account_more_zero',
    'bank_credit', 'critical_account', 'job_skilled_employee', 'male_single',
    'employment_between_1_4_years', 'dwelling_owned', 'employment_less_1_year',
    'existing_credit_paid_back_duly_till_now', 'all_credit_at_this_bank_paid_back_duly_till_now',
    'job_management_officer_etc', 'dwelling_rented', 'dependent']].copy()

X_train, X_test, y_train, y_test = train_test_split(few_variables.drop('credit', axis=1),
    few_variables['credit'], test_size=0.2,
    random_state=42)

clf = MLPClassifier(solver='lbfgs', alpha=1e-5,
    hidden_layer_sizes=(250,), learning_rate='constant', random_state=42)

clf.fit(X_train, y_train)
predictions = clf.predict(X_test)

print(classification_report(y_test, predictions))
print(predictions)

# Score
score = clf.score(X_test, y_test)

# Confusion Matrix
confusion_matrix = confusion_matrix(y_test, predictions)
plt.figure(figsize=(9,9))
sns.heatmap(confusion_matrix, annot=True, fmt=".3f", linewidths=.5, square = True, cmap='Blues')
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
all_sample_title = 'Confusion Matrix -- Accuracy Score: {0}'.format(score)
plt.title(all_sample_title, size = 15);

# ROC -- Test
neural_network_roc_auc = roc_auc_score(y_test, clf.predict(X_test))
fpr, tpr, thresholds = roc_curve(y_test, clf.predict_proba(X_test)[:,1])
plt.figure()
plt.plot(fpr, tpr, label='Neural Network - TEST DATA (area = %0.2f)' % neural_network_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')

```

```

plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Tree_ROC')
plt.show()

# ROC -- Train
neural_network_roc_auc = roc_auc_score(y_train, clf.predict(X_train))
fpr, tpr, thresholds = roc_curve(y_train, clf.predict_proba(X_train)[: ,1])
plt.figure()
plt.plot(fpr, tpr, label='Neural Network - TRAIN DATA (area = %0.2f)' % neural_network_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Tree_ROC')
plt.show()


```

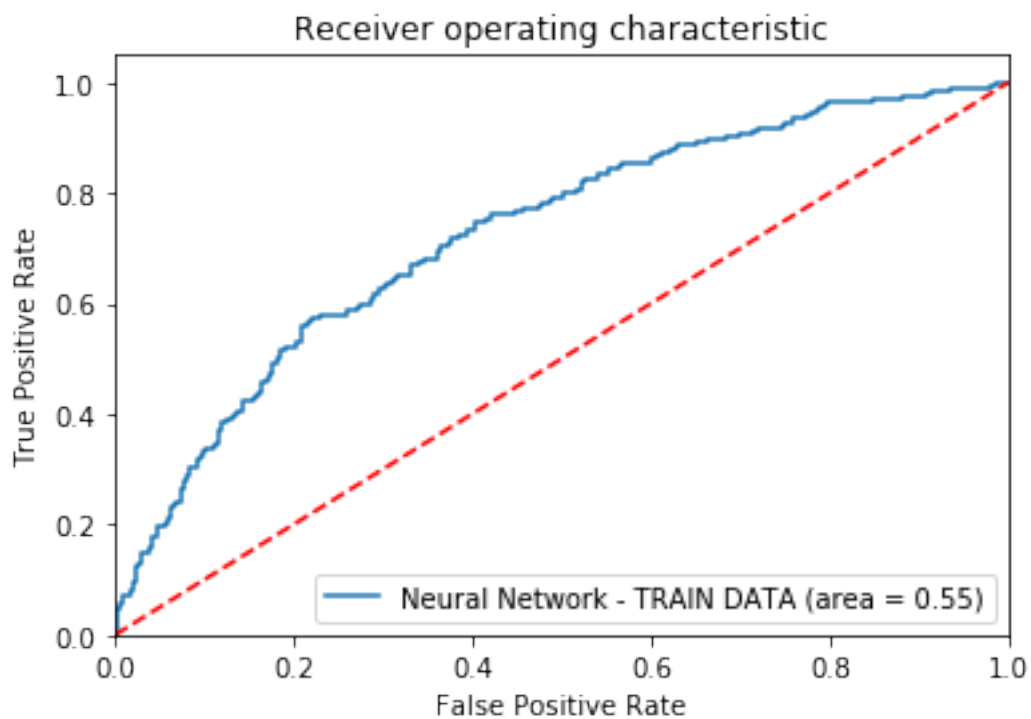
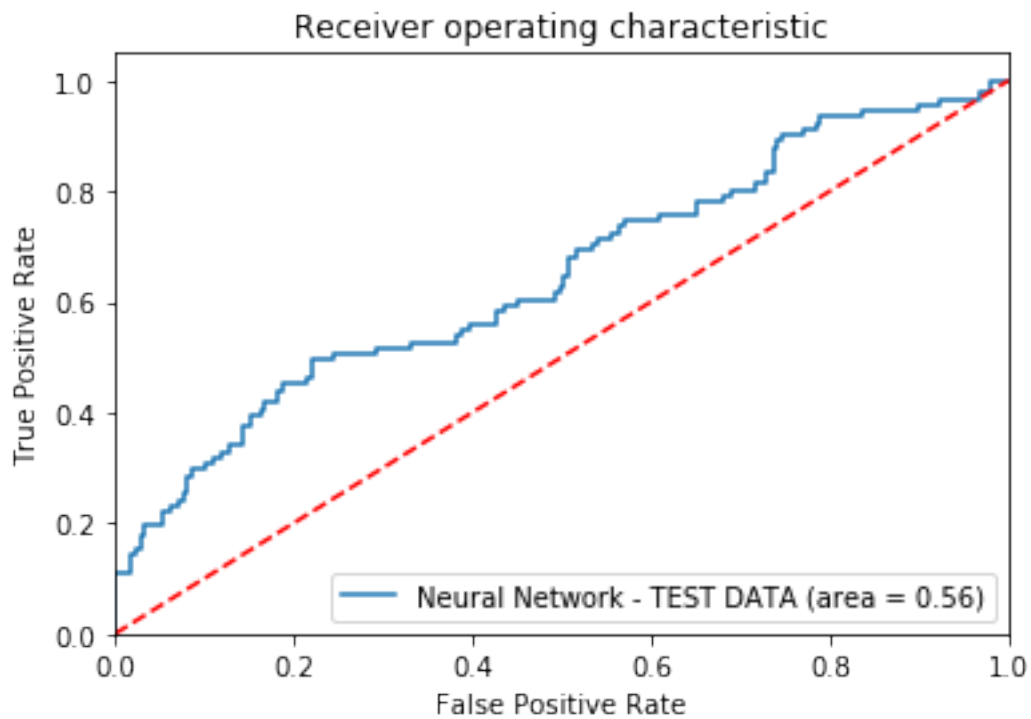
	precision	recall	f1-score	support
0	0.72	0.97	0.83	209
1	0.70	0.15	0.25	91
micro avg	0.72	0.72	0.72	300
macro avg	0.71	0.56	0.54	300
weighted avg	0.72	0.72	0.65	300

```

[0 0 0 0 0 0 0 1 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0
 0 0 1 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0
 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1
 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0
 0 0 0 0]

```





4 Sélection du modèle

Les 3 derniers modèles évalués sont SVM linéaire, SVM radial et le modèle de neurone. ENcore une fois, nous avons jouer avec les paramètres de chacun dans le but d'obtenir le meilleur résultat.

Le modèle retenu est SVM Linéaire et ce parce que son score d'accuracy est le meilleur des 3, à 74%.

Il semble que nos données soient définitivement faites pour être prédites avec un modèle de type linéaire puisque nos deux meilleures estimations proviennent de ce type de modèle.

5 Conclusion

En conclusion, nos résultats sont satisfaisants. Un score autour de 70 et 75% n'est pas mal considérant que les calculs ont été faits sur 1000 échantillons uniquement. Dans tous les cas, ce laboratoire nous a permis d'explorer les différents modèles de prédictions en machine learning. Nous avons pu voir comment explorer les données et par la suite les évaluer en jouant avec les modèles et leurs différents paramètres.

Augustin Commun Chloé Constantineau